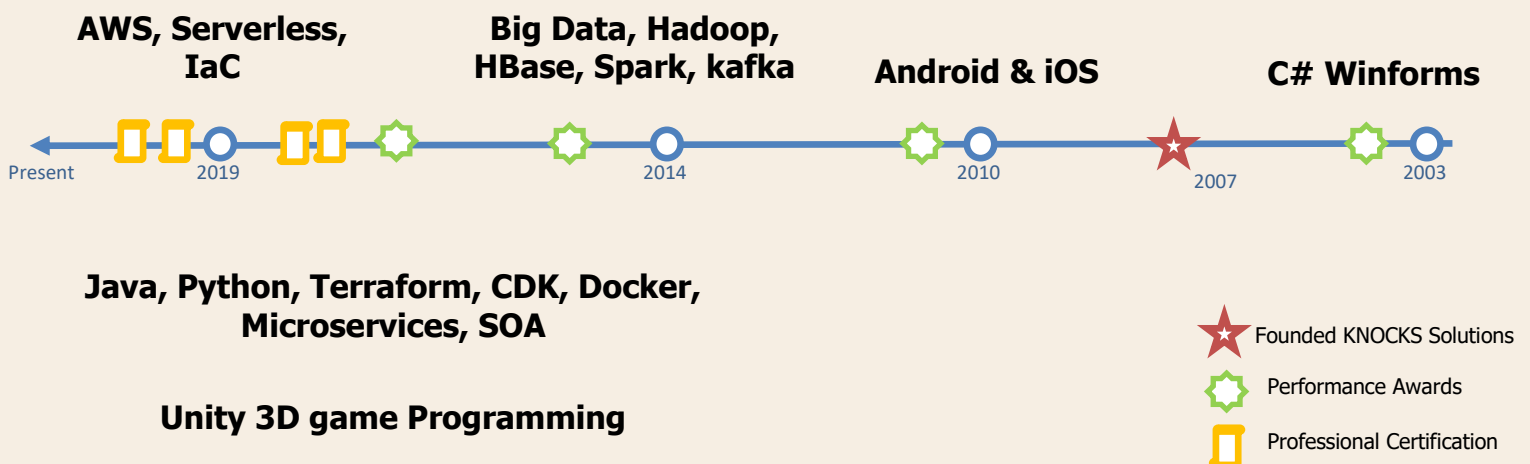


**AWS, Databricks, Confluent Certified Solutions Architect with overall 19+ years of experience, with the last 9+ years in Big Data and Cloud computing.**

**Passionate about programming and design and enjoy working under tight deadlines.**

**Reputed as a problem solver, with ability to work across programming languages and frameworks.**

## Skills Timeline



## About Me

- Technology and programming-language agnostic, curious about emerging trends in computing and big believer in design driven development.
- **Founded and operated KNOCKS Solutions (2007-2010)**, developed third party .NET Winforms Components Suite, like Outlook style Calendar, Visual Studio like Docking Extender, Office style Ribbon Control, Winamp like Skin-extender and various meters and gauges UI components.
- **Self-Learned Unity3D game programming** as a hobby and developed some games for Mobile Platforms hosted at <https://rgyani.com>
- Also dabbled in Machine learning Image recognition projects like License Plate detection, Apparel recommender system using CNNs, YOLO, Colab etc.
- Was involved in various other projects across technologies in previous organizations, and also a lot of POCs and Pitch projects in collaboration with various Sales Team, all of which have been excluded here for the sake of brevity. Also **received a lot of recognitions and awards** throughout my career.

## Abridged Employment History

### Senior Solutions Architect at Luxoft GmbH, Munich

Jan' 2020—Present

- Worked with various clients like Amazon, BMW and Continental on developing their AWS stacks, which involved designing and developing-on various AWS Components, including **Lambdas, Kinesis, API Gateway, Cognito, SNS and SQS, DynamoDB and Neptune DB, Cloudformation, CDK, Greengrass, IoT Core, Glue** etc.
- Built infrastructure pipelines using Terraform and CDK (incl. Custom Resources), Jenkins and AWS CodePipeline.
- Built my reputation as a **Problem-Solver**, implementing some out-of-the-box solutions to solve recurring problems for the clients. Some examples include a custom Grafana plugin to query MapR cluster using Drill REST API (developed in GO and Angular), a custom zeppelin interpreter for git interactions (Java & Angular), a EMR based MapR S3 connector to transfer data from S3 to on-prem MapR cluster (Java, EMR)
- Optimized the analytics pipeline for a client, **bringing the cost down by 90%** in a complex AWS workflow, which was using EFS, AWS Batch, and Managed Apache Airflow with S3.
- Developed a low-cost, vendor neutral solution for a startup, which involved Data parsing via Kafka Streams & Spark from a proprietary format to S3 Parquet, and using Spark or Java or Athena as a Query layer on top (with S3 Select) for use by Machine-Learning and front-end Teams
- Mentored colleagues across various nationalities in AWS, conducting learning and idea sharing sessions, apart from regular code-reviews and collaborative development

### Solutions Architect / Tech Lead at Polaris Wireless, Bangalore

Feb' 2014— Aug' 2019

- Joined Polaris as an iOS/Android Tech lead, but quickly realized their SQL Server based product to store and process the data was not scalable, **Scaled myself into a Big Data Developer**. Built the first POC, sold the idea to Management and Client, even went onsite for first deployment
- Designed and developed a Big Data bare-metal as well as cloud (AWS) based solution, to store and profile mobile locations of all subscribers of a country for use by investigative agencies.
- Worked on various NoSQL offerings like **HBase, Cassandra, MongoDB, Solr, Kafka, Spark**. Using **Java, Scala and Python** developed the full product, including KPIs, benchmarked ingestion and performance; lead a team of 9 full time engineers.
- Also involved in designing various microservices to generate profile, algorithms for comparison of the profiles (in near-real time) to find familial, communication and location patterns between subscribers (Named on Patent)
- All these use HBase over Hadoop for storage and are optimized by Distributed Computing including **Map Reduce** and **Spark**. **Kafka** is used as the communication channel between the microservices.
- Developed the next-gen of the product, where-in Profiling is done in-line with ingestion, with **AWS S3** as the storage medium and **AWS EMR** and **AWS MSK** for ingestion with some custom cardinality solutions using Count-Min Sketch, BloomFilters, HyperLogLog etc

## Education and Professional Certifications

**B.E. from Netaji Subhas Institute of Technology, Delhi** April 1999 – June 2003

**AWS Certified Solutions Architect - Professional** SAP-C01

**AWS Certified Solutions Architect - Associate** SAA-C02

**Databricks Certified Developer: Apache Spark 2.X** 2019

**Confluent Certified Developer for Apache Kafka** 2019